

1 **Supplementary**

2 This is the supplementary material to the paper, Adapting by Analogy: OOD Generalization of
3 Visuomotor Policies via Functional Correspondence. More details can be found on our project page
4 <https://anon-corl2025.github.io/project-page/>

5 **1 Hardware Experiment Setup**



Figure 1: Our hardware experiment setup, we use a Franka Research 3 robot, with a UMI gripper. The RealSense D435 wrist camera, and Zed mini 2i third person camera are placed as shown.

6 **2 Tasks**

7 We conduct our experiments on two real-world tasks. The first task is **sweep-trash**, wherein robot
8 must sweep trash towards different goals, based on whether the trash is organic and recycling. For
9 evaluation, we divide the task in two sub-goals (A) properly aligning the wiper with the trash, (B)
10 sweeping to the correct location. The next task is **object-in-cup**, where-in a robot arm is tasked with
11 picking up a object such as a marker or a pen and dropping it in a mug. Markers—which are grasped
12 above their center-of-mass—need to be dropped into the mug from the bottom, and pens—which
13 are grasped below their center-of-mass—need to be dropped from the front. We divide the task in
14 3 sub-goals (A) grasping the object, (B) picking the correct behavior mode based on the grasp, and
15 (C) dropping object into the cup. Fig. 2 demonstrates the various modes and the sub-goals for the
16 two tasks.

17 We evaluate both tasks on the in-distribution conditions and OOD conditions induced by back-
18 ground and novel objects. The OOD environments are shown in Fig. 3 For **sweep-trash** our
19 ID environments are $E_{ID\text{-trash}} := \{E_{ID}^{paper}, E_{ID}^{M\&Ms}\}$. The OOD environments are $E_{OOD\text{-Trash}} :=$
20 $\{E_{OOD}^{doritos}, E_{OOD}^{napkin}, E_{OOD}^{thumb-tack}, E_{OOD}^{paper-bg}, E_{OOD}^{M\&M-bg}\}$.




Figure 2: The training demonstrations for our two tasks, with their sub-goals(A, B, C). For the object in cup task, the pen is grasped below the center-of-mass, and is dropped into the mug from the front. The marker is grasped above the center-of-mass and is dropped into the mug from the bottom. For the sweep trash task, paper (i.e., recycling) is swept up, and M\$Ms (i.e., organic) is swept down.

21 For **object-in-cup** our ID environments are $E_{\text{ID-object}} := \{E_{\text{ID}}^{\text{marker}}, E_{\text{ID}}^{\text{pen}}\}$. The OOD environments
 22 are $E_{\text{OOD-object}} := \{E_{\text{OOD}}^{\text{pencil}}, E_{\text{OOD}}^{\text{battery}}, E_{\text{OOD}}^{\text{block}}, E_{\text{OOD}}^{\text{marker-bg}}, E_{\text{OOD}}^{\text{pen-bg}}\}$.




Figure 3: Our OOD environments for both the sweep-trash and object-in-cup task

23 3 Additional Results

24 3.1 How much does **ABA** improve the policy’s sub-goal level closed-loop performance?

25 Fig. 4 shows that on the sweep-trash task, both **ABA** and **Vanilla** are able to successfully accomplish
 26 both subgoals on $E_{\text{ID}}^{M\&Ms}$. However, in $E_{\text{ID}}^{\text{paper}}$, **Vanilla** fails at aligning the wiper with the paper
 27 trash (subgoal A) 10% of the times, and sweeps paper incorrectly (subgoal B) 30% of the times.
 28 **ABA** maintains 100% performance on $E_{\text{ID}}^{\text{paper}}$.

29 Showing a similar trend, both **Vanilla** and **ABA** show 100% success rate on the $E_{\text{ID}}^{\text{marker}}$ for the
 30 object-in-cup task. On the $E_{\text{ID}}^{\text{pen}}$, while both **Vanilla** and **ABA** are able to grasp the pen (subgoal
 31 A) 100% of the times, **ABA** improves over **Vanilla** by 40% at picking the right mode (subgoal B),
 32 showing a 100% success rate. Finally, since dropping the pen into the cup (subgoal C) is the most
 33 fine-grained aspect of the task, both **ABA** and **Vanilla** struggle but **ABA** still improves over **Vanilla**
 34 by 50%.




Figure 4: Subgoal Success in each ID Environment. We report the subgoal level task success rate averaged across 10 rollouts. For both the sweep-trash and the object-in-cup tasks, we see that **ABA** consistently achieves the highest task success rate compared to baselines.




Figure 5: Subgoal Success in each OOD Environment, induced by changing the background. The success rate is averaged across 10 rollouts. **ABA** again consistently achieves the highest task success rate compared to baselines.

35 Next, we compare **ABA** and **Vanilla** across OOD environments, induced using a novel background
 36 ($E_{OOD}^{paper-bg}$, $E_{OOD}^{M\&Ms-bg}$, E_{OOD}^{pen-bg} , $E_{OOD}^{marker-bg}$).

37 Fig. 5 shows that on the sweep-trash task the performance trends are similar to the ID environments,
 38 although interestingly instead of $E_{OOD}^{paper-bg}$, **Vanilla** now shows poorer performance on subgoal B
 39 of $E_{OOD}^{M\&Ms-bg}$. **ABA** shows 100% success rate on both goals of both $E_{OOD}^{M\&Ms-bg}$ and $E_{OOD}^{paper-bg}$. On the
 40 object-in-cup task, **Vanilla** struggles on all subgoals of both E_{OOD}^{pen-bg} , $E_{OOD}^{marker-bg}$ environments. **ABA**




Figure 6: Subgoal Success in each OOD Environment with 3 novel objects for both sweep-trash and object-in-cup task. The success rate is averaged across 10 rollouts. **ABA** again consistently achieves the highest task success rate compared to baselines.

41 improves over **Vanilla** on both environments showing a 100% performance on all subgoals of both
 42 environments, except subgoal C of the $E_{\text{OOD}}^{\text{pen-bg}}$, where it shows an 80% success rate. This shows
 43 that a with novel background, **Vanilla** fails to even grasp the objects, however interventions with
 44 ID observations ignores the OOD conditions induced by the novel background, allowing **ABA** to
 45 uphold closed loop performance under the OOD environments.

46 Finally, on OOD environments induced by novel object categories for the sweep-trash task we ob-
 47 serve from Fig. 6 that while **Vanilla** is able to align the wiper with the trash, it fails to pick the correct
 48 direction for sweeping the trash (subgoal B), as visual features are not enough to decide whether the
 49 trash is organic or recycling. **ABA** is able to successfully accomplish both subgoals for all novel
 50 objects as the relevant features for deciding the trash type are supplied by the expert as functional
 51 correspondences.

52 Since the object-in-cup task is more challenging, **Vanilla** is only performant at grasping (subgoal
 53 A). It is able to grasp the pencil with 40%, the battery with 90%, and the jenga-block with 80%
 54 success-rate. However, the sizes of the objects are such that they can only be dropped into the mug
 55 from the top (subgoal B), however **Vanilla** is not able to infer these features solely from the training
 56 data and hence fails at subgoal B and C. With **ABA**, the expert language feedback helps establish
 57 the correct functional correspondences, leading to an improvement in the performance across all
 58 subgoals.

59 3.2 What kind of features maximally improve the sub-goal level performance for 60 observation interventions based methods?

61 As shown in Fig. 4, all intervention based method demonstrate a 100% task success on all sub-
 62 goals of the sweep-trash task, in the ID environments. On the object-in-cup task intervention based
 63 methods again perform comparably on the $E_{\text{ID}}^{\text{marker}}$, however on the $E_{\text{ID}}^{\text{pen}}$ both **PolicyEmbed** and
 64 **DINOEmbed** perform worse as compared to **ABA** on all subgoals.

65 As shown in Fig. 5, under a novel background, intervention based methods perform comparably
 66 on the sweep-trash task. On the object-in-cup task, **PolicyEmbed** performs worse compared to
 67 both **DINOEmbed** and **ABA**, whereas **DINOEmbed** performs comparably with **ABA**. This can be
 68 attributed to the ability of dino features to perform dense correspondence matching, specially across
 69 objects in the same semantic class.

70 Fig. 6 shows that under novel objects both **PolicyEmbed** and **DINOEmbed** struggle. Since **Pol-**
 71 **icyEmbed** relies on the policy embeddings, under ‘doritos’ and ‘napkin’ it sweeps them in either
 72 direction. **DINOEmbed** matches the visual features and since napkin closely resembles paper, it is
 73 able to correctly sweep napkin as recycling, and fails on other objects.

74 For the object-in-cup task, because policy embeddings and visual features alone are not enough to
75 match the objects with the ID sample that lead to the desired behavior mode, both **PolicyEmbed**
76 and **DINOEmbed** perform worse as compared to **ABA**.